

Run: 2026-02-18 · 14:42 MYT Operator: Tuan Source: the console above

Batch report

ORD-1279 → ORD-1284

Six orders through one queue. The console above routed them live; this is the receipt it printed — same lane names, same 6 orders, same single failover (ORD-1281, 429 → retried on another lane). Anyone can read it; engineers can verify it.

THROUGHPUT

6 orders

1 batch · 6 lanes

P95 TIME

298 ms

fastest 184 ms · slowest 388 ms

FAILOVER

1 · recovered

ORD-1281 → retried · 200

SUCCESS

100 %

6/6 delivered

WHAT THIS PROVES

One queue handles many lanes, and a hiccup doesn't lose the order. The report mirrors the console's live state so a non-technical reader follows the story and a technical reader can audit it.

ORDER	LANE	MODEL (UNDER THE HOOD)	TIME	RESULT
ORD-1284	Express lane	kimi-k3-max · openai-custom	184 ms	200 · DELIVERED
ORD-1283	Batch lane	deepseek-reasoner · deepseek-alias	262 ms	200 · DELIVERED
ORD-1282	Priority lane	claude-sonnet-4 · anthropic-native	210 ms	200 · DELIVERED
ORD-1281	Standard lane → Priority lane	gpt-4.1-mini → claude-sonnet-4	298 ms → 210 ms	429 → 200 · FAILOVER
ORD-1280	Express lane	kimi-k3-max · openai-custom	184 ms	200 · DELIVERED
ORD-1279	Coder lane	qwen2.5-coder · dashscope	224 ms	200 · DELIVERED

In plain words: think of lanes as checkout counters. Six orders came in one line; the system sent each to the best open counter. One counter was busy (429) — the order stepped to the next counter and went through. No order lost.

LANE	MODEL	AVG TIME	HEALTH	SHARE
Express lane	kimi-k3-max	184 ms	OK	dedicated
Batch lane	deepseek-reasoner	262 ms	OK	shared
Priority lane	claude-sonnet-4	210 ms	OK	dedicated
Standard lane	gpt-4.1-mini	298 ms	OK	shared
Overflow lane	grok-4	388 ms	WARM	dedicated
Coder lane	qwen2.5-coder	224 ms	OK	shared

ROUTING SIGNALS — SAME 3 CHARTS AS L2

Time (ms), volume, and cost over the last 14 hours — the console's three sparklines, frozen here for the record.



These are the console's **golden signals** relabelled in plain words — time = how fast, volume = how many, cost = how much. Engineers still see latency/tokens/cost underneath; everyone else sees speed/scale/spend.

WHY IT MATTERS (FOR EVERYONE)

A queue that can reroute keeps promises. Customers see orders arrive; finance sees cost stay flat; engineers see p95 and failover rate. One batch, one receipt, many readers.

Every order, start to finish.

The complete run in arrival order — queue time, lane assignment, retries, and final status for each order. This is the same data the console streams, frozen at batch close.

ORDER	QUEUED	LANE	QUEUE → DONE	RETRIES	RESULT
ORD-1279	14:42:01	Coder lane	0 ms → 224 ms	0	200
ORD-1280	14:42:01	Express lane	2 ms → 184 ms	0	200
ORD-1281	14:42:02	Standard → Priority	298 ms → 210 ms	1	200 • RETRIED
ORD-1282	14:42:02	Priority lane	1 ms → 210 ms	0	200
ORD-1283	14:42:03	Batch lane	3 ms → 262 ms	0	200
ORD-1284	14:42:03	Express lane	0 ms → 184 ms	0	200

Reading this table: "queued" is when the order entered the system. Queue time under 5 ms means a lane was free immediately — only ORD-1281 waited (its first lane returned 429 busy). Total added latency from the retry: 88 ms.

How the lanes held up.

Per-lane utilization and error rates across the full 24-hour window ending with this batch. The Standard lane's 429 was load-related — its recovery is tracked below.

LANE	LOAD (THIS BATCH)	ERROR RATE • 24H	P95 • 24H	STATUS
Express lane		0.00%	191 ms	HEALTHY
Priority lane		0.01%	224 ms	HEALTHY
Batch lane		0.02%	270 ms	HEALTHY
Standard lane		2.10%	305 ms	RECOVERING
Overflow lane		0.00%	392 ms	STANDBY
Coder lane		0.01%	231 ms	HEALTHY

Capacity note: Standard lane ran at 82% of its safe concurrency when ORD-1281 arrived — the 429 was the guardrail working as designed, not an outage. Load shed to 61% within 90 seconds; no other order touched that lane during the window.

FAILOVER PATH TAKEN

Standard lane (429) → router cooldown 88 ms → Priority lane → delivered.
The Overflow lane stayed in reserve the entire batch.

One hiccup. No lost order.

ORD-1281 hit a busy Standard lane (429). The router cooled down for 88 ms, retried on Priority lane (claude-sonnet-4), and delivered in 210 ms. This is the same event the console highlighted live — match the ID in both places.

STEP	WHAT HAPPENED	TIME
1	ORD-1281 routed to Standard lane (gpt-4.1-mini) — busy	298 ms · 429
2	Router cooldown · lane health re-check	88 ms
3	Retry → Priority lane (claude-sonnet-4)	210 ms
4	Delivered · no customer-visible retry	200

OUTCOME

Batch success 6/6. Effective p95 unchanged at 298 ms. Cost delta +\$0.002 for the retry — within noise. Overflow lane was never needed.

Audit link: verify ORD-1281 in the console's "Live fulfilment" list — it pulses with a tint on load. The report's row uses the identical ID and route. Same data, two views.

What it cost. Who checked it.

Spend for this batch against the daily budget, plus the sign-off chain. The retry's cost delta is visible and immaterial — that's the point.

TOKEN SPEND

\$0.847 of \$40 budget

2.1% of daily

RETRY DELTA

+\$0.002 failover cost

within noise

COST / ORDER

\$0.141 avg

↓ flat vs 7d avg

AUDIT STEP	ACTOR	TIMESTAMP	HASH
Batch submitted	ops-console	14:42:01.204	b7e2...41ac
Lanes assigned	model-router v2.4	14:42:01.208	c91d...08f2
Failover executed	model-router v2.4	14:42:02.511	a3f8...77b9
Batch closed · report emitted	report-writer	14:42:04.967	e55c...d210
Signed off	Tuan · Systems	14:44:12.000	—

PREPARED BY	DATE	RETENTION
TUAN · SYSTEMS	2026-02-18	90 DAYS