

Run: 2026-02-18 · 14:42 MYT Operator: Tuan Source: the console above

Batch report

ORD-1279 → ORD-1284

Six orders through one queue. The console above routed them live; this is the receipt it printed — same lane names, same 6 orders, same single failover (ORD-1281, 429 → retried on another lane). Anyone can read it; engineers can verify it.

THROUGHPUT

6 orders

1 batch · 6 lanes

P95 TIME

298 ms

fastest 184 ms · slowest 388 ms

FAILOVER

1 · recovered

ORD-1281 → retried · 200

SUCCESS

100 %

6/6 delivered

WHAT THIS PROVES

One queue handles many lanes, and a hiccup doesn't lose the order. The report mirrors the console's live state so a non-technical reader follows the story and a technical reader can audit it.


ORDER	LANE	MODEL (UNDER THE HOOD)	TIME	RESULT
ORD-1284	Express lane	kimik3max · openai-custom	184 ms	200 · DELIVERED
ORD-1283	Batch lane	deepseek-reasoner · deepseek-alias	262 ms	200 · DELIVERED
ORD-1282	Priority lane	claude-sonnet-4 · anthropic-native	210 ms	200 · DELIVERED
ORD-1281	Standard lane → Priority lane	gpt-4.1-mini → claude-sonnet-4	298 ms → 210 ms	429 → 200 · FAILOVER
ORD-1280	Express lane	kimik3max · openai-custom	184 ms	200 · DELIVERED
ORD-1279	Coder lane	qwen2.5-coder · dashscope	224 ms	200 · DELIVERED

In plain words: think of lanes as checkout counters. Six orders came in one line; the system sent each to the best open counter. One counter was busy (429) — the order stepped to the next counter and went through. No order lost.

LANE	MODEL	AVG TIME	HEALTH	SHARE
Express lane	kimi-k3-max	184 ms	OK	dedicated
Batch lane	deepseek-reasoner	262 ms	OK	shared
Priority lane	claude-sonnet-4	210 ms	OK	dedicated
Standard lane	gpt-4.1-mini	298 ms	OK	shared
Overflow lane	grok-4	388 ms	WARM	dedicated
Coder lane	qwen2.5-coder	224 ms	OK	shared

ROUTING SIGNALS — SAME 3 CHARTS AS L2

Time (ms), volume, and cost over the last 14 hours — the console's three sparklines, frozen here for the record.



These are the console's **golden signals** relabelled in plain words — time = how fast, volume = how many, cost = how much. Engineers still see latency/tokens/cost underneath; everyone else sees speed/scale/spend.

WHY IT MATTERS (FOR EVERYONE)

A queue that can reroute keeps promises. Customers see orders arrive; finance sees cost stay flat; engineers see p95 and failover rate. One batch, one receipt, many readers.

One hiccup. No lost order.

ORD-1281 hit a busy lane (429). The router retried on Priority lane (claude-sonnet-4) and delivered in 210 ms. The report captures the same event the console highlighted live — match the ID in both places.

STEP	WHAT HAPPENED	TIME
1	ORD-1281 routed to Standard lane (gpt-4.1-mini) — busy	298 ms · 429
2	Router retried → Priority lane (claude-sonnet-4)	210 ms
3	Delivered · no retry needed by the customer	200

OUTCOME

Batch success 6/6. Effective p95 unchanged. Cost delta +\$0.002 for the retry — within noise. The same ORD-1281 row is highlighted live in the console above; the report freezes it for audit.

Audit link: verify ORD-1281 in the console's "Live fulfilment" list — it pulses with a tint on load. The report's row uses the identical ID and route. Same data, two views.

PREPARED BY	DATE	SOURCE
TUAN · SYSTEMS	2026-02-18	L2 CONSOLE (THIS PAGE'S SOURCE)

Built as a portfolio proof: the console is interactive; the report is the artefact it emits. Both are hand-built — no external BI tool — so the craft is the message.